

```
In [1]: # This model is using Linear Regression to predict house price based on the accessibility to the nearest MRT station.

#installing python libraries.

import pandas as pd
import numpy as np
import seaborn as sns
%matplotlib inline
```

```
In [2]: pwd
```

```
Out[2]: 'C:\\Users\\rehnu'
```

```
In [3]: cd C:\\Users\\rehnu\\OneDrive\\Desktop\\Python!\\Datasets for Kaggle Projects
C:\\Users\\rehnu\\OneDrive\\Desktop\\Python!\\Datasets for Kaggle Projects
```

```
In [4]: Housing_data = pd.read_excel("Real_estate_valuation_data set.xlsx")
```

```
In [5]: #Showing first few rows of data.
Housing_data.head()
```

```
Out[5]:
```

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
0	1	2012.916667	32.0	84.87882	10	24.98298	121.54024	37.9
1	2	2012.916667	19.5	306.59470	9	24.98034	121.53951	42.2
2	3	2013.583333	13.3	561.98450	5	24.98746	121.54391	47.3
3	4	2013.500000	13.3	561.98450	5	24.98746	121.54391	54.8
4	5	2012.833333	5.0	390.56840	5	24.97937	121.54245	43.1

In [6]: *#showing last few rows of data.*
Housing_data.tail()

Out[6]:

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
409	410	2013.000000	13.7	4082.01500	0	24.94155	121.50381	15.4
410	411	2012.666667	5.6	90.45606	9	24.97433	121.54310	50.0
411	412	2013.250000	18.8	390.96960	7	24.97923	121.53986	40.6
412	413	2013.000000	8.1	104.81010	5	24.96674	121.54067	52.5
413	414	2013.500000	6.5	90.45606	9	24.97433	121.54310	63.9

In [7]: *#Checking any null values.*
Housing_data.isnull().sum()

Out[7]: No 0
X1 transaction date 0
X2 house age 0
X3 distance to the nearest MRT station 0
X4 number of convenience stores 0
X5 latitude 0
X6 longitude 0
Y house price of unit area 0
dtype: int64

In [8]: *#Showing data types of all columns.*
Housing_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 414 entries, 0 to 413
Data columns (total 8 columns):
No                414 non-null int64
X1 transaction date    414 non-null float64
X2 house age          414 non-null float64
X3 distance to the nearest MRT station  414 non-null float64
X4 number of convenience stores    414 non-null int64
X5 latitude           414 non-null float64
X6 longitude          414 non-null float64
Y house price of unit area    414 non-null float64
dtypes: float64(6), int64(2)
memory usage: 26.0 KB
```


In [9]: *#Showing the total number of rows and columns in the dataset.*
Housing_data.shape

Out[9]: (414, 8)

In [10]: *#Showing the scatter plot of "distance to the nearest MRT station" vs "house price".*

```
sns.pairplot(Housing_data, x_vars = ['X3 distance to the nearest MRT station'], y_vars = ['Y house price of unit area'], height = 10, aspect = 1.0, kind = 'scatter')
```

Out[10]: <seaborn.axisgrid.PairGrid at 0x1b7dcef8088>



In [11]: *#Declaring the independent variable.*

```
X = Housing_data['X3 distance to the nearest MRT station']
X.head()
```

Out[11]:

0	84.87882
1	306.59470
2	561.98450
3	561.98450
4	390.56840

Name: X3 distance to the nearest MRT station, dtype: float64

In [12]: *#Declaring the dependent variable.*

```
Y = Housing_data['Y house price of unit area']
Y.head()
```

Out[12]:

0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

Name: Y house price of unit area, dtype: float64

In [13]: *#Splitting the dataset into train and test.*

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, train_size = 0.7, random_state = 100)
```

In [14]: *#Showing the train data.*

```
print(X_train.shape)
```

(289,)

In [15]: *#Fixing the column issues in the train and test dataset as previously they showed only the number of rows.*

```
import numpy as np
```

```
X_train = X_train[:, np.newaxis]
X_test = X_test[:, np.newaxis]
```

In [16]:

```
print(X_train.shape)
```

(289, 1)

```
In [17]: #Fixing linear regression model.

from sklearn.linear_model import LinearRegression

lr = LinearRegression()

lr.fit (X_train, Y_train)
```

```
Out[17]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [18]: #Printing the intercept and slope of the linear regression model.

print (lr.intercept_)
print(lr.coef_)

45.825776046478694
[-0.00714734]
```

```
In [19]: #Declaring a prediction variable to test the model.

Y_pred = lr.predict(X_test)
```

```
In [20]: #Showing prediction and test variable size.


print(Y_pred.shape)
print(Y_test.shape)

(125,)
(125,)
```

In [21]: *#Showing comparison between actual and predicted values.*

```
import matplotlib.pyplot as plt
c = [i for i in range (0,125,1)]
fig=plt.figure()
plt.plot(c,Y_test, color="blue", linewidth=2.5, linestyle="-")
plt.plot(c,Y_pred, color="red", linewidth=2.5, linestyle="-")
fig.suptitle('Actual and Predicted', fontsize=20)
plt.xlabel('Index', fontsize=18)
plt.ylabel('Unit price',fontsize=16)
```


Out[21]: Text(0, 0.5, 'Unit price')



In [22]: *#Showing the residual plot.*

```
c = [i for i in range (0,125,1)]
fig=plt.figure()
plt.plot(c,Y_test-Y_pred, color="blue", linewidth=2.5, linestyle="-")
fig.suptitle('Error Terms', fontsize=20)
plt.xlabel('Index', fontsize=18)
plt.ylabel('Y_test-Y_pred',fontsize=16)
```

Out[22]: Text(0, 0.5, 'Y_test-Y_pred')



In [23]: *#Calculating mean squared error.*

```
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error (Y_test, Y_pred)
```

In [24]: *#Calculating the R-squared value.*

```
r_squared=r2_score (Y_test, Y_pred)
```

In [25]:


```
print('Mean_Squared_Error :', mse)
print('r_squared_value :',r_squared)
```

```
Mean_Squared_Error : 68.6178057851028
r_squared_value : 0.5104359968219091
```

In [26]: *#Showing scatter plot of actual and predicted values.*

```
import matplotlib.pyplot as plt
plt.scatter(Y_test, Y_pred)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
```

Out[26]: Text(0, 0.5, 'Predicted Y')



In [27]: *# The Simple Linear Regression exercise has an R-squared value of 0.51 and Mean Squared Error of 68.61, which indicate that the relationship between the two variables do not have a very strong Linear relationship. That is why the actual and predicted house unit price differed so much from each other. A non-linear regression model could be tried in this case to come up with a better model.*

In []: